

Complete VectorStore Setup Guide

Step 1: Create PostgreSQL Database and User

1.1 Connect to PostgreSQL as superuser

```
bash  
  
sudo -u postgres psql
```

1.2 Create database and user (run these SQL commands in PostgreSQL)

```
sql  
  
-- Create a new database  
CREATE DATABASE vector_db;  
  
-- Create a new user  
CREATE USER vector_user WITH PASSWORD 'your_secure_password';  
  
-- Grant privileges to the user  
GRANT ALL PRIVILEGES ON DATABASE vector_db TO vector_user;  
  
-- Connect to the new database  
\c vector_db  
  
-- Create required extensions  
CREATE EXTENSION IF NOT EXISTS vector;  
CREATE EXTENSION IF NOT EXISTS timescaledb;  
  
-- Grant usage on schema to the user  
GRANT USAGE ON SCHEMA public TO vector_user;  
GRANT CREATE ON SCHEMA public TO vector_user;  
GRANT ALL PRIVILEGES ON ALL TABLES IN SCHEMA public TO vector_user;  
GRANT ALL PRIVILEGES ON ALL SEQUENCES IN SCHEMA public TO vector_user;  
  
-- Exit PostgreSQL  
\q
```

Step 2: Test Database Connection

2.1 Test the connection with your new credentials

```
bash  
  
psql -h localhost -p 5432 -U vector_user -d vector_db
```

2.2 Verify extensions are installed

```
sql

SELECT extname, extversion FROM pg_extension WHERE extname IN ('vector', 'timescaledb');

\q
```

Step 3: Set Environment Variables

```
bash

# Set your Google API key
export GOOGLE_API_KEY="your_actual_gemini_api_key_here"

# Optional: Set database URL as environment variable
export DATABASE_URL="postgresql://vector_user:your_secure_password@localhost:5432/vector_db"
```

Step 4: Install Required Python Packages

```
bash

pip install timescale-vector google-generativeai python-dotenv psycpg2-binary
```

Step 5: Fixed Python Code

Save this as `vector_store_fixed.py`:


```

import google.generativeai as genai
from typing import List, Dict, Any
import logging
import time
import os
from datetime import timedelta
import psycopg2
from urllib.parse import urlparse

# Try to import the vector database client
try:
    from timescale_vector import client
    VECTOR_CLIENT_AVAILABLE = True
    print("✅ Timescale Vector client imported successfully")
except ImportError as e:
    VECTOR_CLIENT_AVAILABLE = False
    print(f"❌ Timescale Vector not available: {e}")
    print("Install with: pip install timescale-vector")

def get_settings():
    """Get database settings from environment or use defaults"""

    # Get from environment variable or use default
    database_url = os.getenv("DATABASE_URL", "postgresql://vector_user:your_secure_password@loc

    return {
        "database": {
            "service_url": database_url
        },
        "vector_store": {
            "table_name": "vector_embeddings",
            "embedding_dimensions": 768,
            "time_partition_interval": timedelta(days=7) # Use timedelta object instead of str
        }
    }

def test_database_connection(database_url: str) -> bool:
    """Test if we can connect to the database"""
    try:
        # Parse the database URL
        parsed = urlparse(database_url)

        # Connect to the database
        conn = psycopg2.connect(

```

```

        host=parsed.hostname,
        port=parsed.port or 5432,
        user=parsed.username,
        password=parsed.password,
        database=parsed.path[1:] # Remove Leading slash
    )

    # Test the connection
    cursor = conn.cursor()
    cursor.execute("SELECT version();")
    version = cursor.fetchone()
    print(f"✅ Database connection successful: {version[0][:50]}...")

    # Check if required extensions exist
    cursor.execute("SELECT extname FROM pg_extension WHERE extname IN ('vector', 'timescale")
    extensions = cursor.fetchall()
    ext_names = [ext[0] for ext in extensions]

    if 'vector' in ext_names:
        print("✅ Vector extension is installed")
    else:
        print("❌ Vector extension is NOT installed")

    if 'timescaledb' in ext_names:
        print("✅ TimescaleDB extension is installed")
    else:
        print("❌ TimescaleDB extension is NOT installed")

    cursor.close()
    conn.close()
    return True

except Exception as e:
    print(f"❌ Database connection failed: {e}")
    return False

```

```

class VectorStore:
    """A class for managing vector operations and database interactions."""

    def __init__(self):
        """Initialize settings and Gemini embedding client."""
        self.settings = get_settings()
        self.vector_settings = self.settings["vector_store"]

        # Gemini embedding model
        self.embedding_model = "models/embedding-001"

```

```

# Set Gemini API key
api_key = os.getenv("GOOGLE_API_KEY")
if not api_key:
    raise ValueError("GOOGLE_API_KEY environment variable is not set")

genai.configure(api_key=api_key)
print("✅ Gemini API configured")

# Test database connection first
db_url = self.settings["database"]["service_url"]
if not test_database_connection(db_url):
    print("❌ Database connection test failed")
    self.vec_client = None
    return

# Initialize vector database client
if VECTOR_CLIENT_AVAILABLE:
    try:
        print(f"🔄 Initializing vector client...")
        print(f"🔄 Database URL: {db_url}")
        print(f"🔄 Table name: {self.vector_settings['table_name']}")
        print(f"🔄 Embedding dimensions: {self.vector_settings['embedding_dimensions']}")
        print(f"🔄 Time partition: {self.vector_settings['time_partition_interval']}")

        self.vec_client = client.Sync(
            service_url=db_url,
            table_name=self.vector_settings["table_name"],
            num_dimensions=self.vector_settings["embedding_dimensions"],
            time_partition_interval=self.vector_settings["time_partition_interval"],
        )
        print("✅ Vector database client initialized successfully")

    except Exception as e:
        print(f"❌ Failed to initialize vector database client: {e}")
        print(f"❌ Error type: {type(e).__name__}")
        print(f"❌ Error details: {str(e)}")
        self.vec_client = None
    else:
        self.vec_client = None
        print("⚠️ Vector database client not available.")

def get_embedding(self, text: str) -> List[float]:
    """Generate embedding for the given text using Gemini."""
    text = text.replace("\n", " ").strip()
    if not text:
        return []

```

```

start_time = time.time()

try:
    response = genai.embed_content(
        model=self.embedding_model,
        content=text,
        task_type="retrieval_document"
    )
    embedding = response["embedding"]

    elapsed_time = time.time() - start_time
    print(f"✅ Embedding generated in {elapsed_time:.3f} seconds")

    return embedding
except Exception as e:
    print(f"❌ Failed to get embedding: {e}")
    return []

def create_table(self) -> bool:
    """Create the vector table"""
    if not self.vec_client:
        print("❌ Vector database client not available")
        return False

    try:
        print("🔄 Creating vector table...")
        self.vec_client.create_tables()
        print("✅ Vector table created successfully")
        return True
    except Exception as e:
        print(f"❌ Failed to create table: {e}")
        return False

def insert_vector(self, text: str, metadata: Dict[str, Any] = None, doc_id: str = None) ->
    """Insert a vector into the database."""
    if not self.vec_client:
        print("❌ Vector database client not available")
        return False

    print(f"🔄 Processing text: '{text[:50]}...'")
    embedding = self.get_embedding(text)
    if not embedding:
        print("❌ Failed to generate embedding")
        return False

    try:

```

```

# Prepare the record
record_id = doc_id or f"doc_{int(time.time() * 1000)}"

# Create record in the format expected by timescale-vector
record = {
    "id": record_id,
    "metadata": metadata or {},
    "contents": text,
    "embedding": embedding
}

print(f"📁 Inserting record with ID: {record_id}")

# Insert into the vector database
self.vec_client.upsert([record])
print(f"✅ Successfully inserted vector")
return True

except Exception as e:
    print(f"❌ Failed to insert vector: {e}")
    print(f"❌ Error type: {type(e).__name__}")
    return False

def search_similar(self, query_text: str, limit: int = 5) -> List[Dict]:
    """Search for similar vectors in the database."""
    if not self.vec_client:
        print("❌ Vector database client not available")
        return []

    print(f"🔍 Searching for: '{query_text}'")
    query_embedding = self.get_embedding(query_text)
    if not query_embedding:
        print("❌ Failed to generate query embedding")
        return []

    try:
        results = self.vec_client.search(
            query_embedding,
            limit=limit
        )
        print(f"✅ Found {len(results)} similar documents")
        return results
    except Exception as e:
        print(f"❌ Search failed: {e}")
        return []

def test_functionality(self):

```



```

"""Test the complete functionality"""
print("\n🧪 Testing VectorStore functionality...")

# Test 1: Embedding generation
test_text = "Machine learning is transforming how we process data."
print(f"\n1. Testing embedding generation...")
embedding = self.get_embedding(test_text)

if not embedding:
    print("❌ Embedding test failed")
    return

print(f"✅ Generated {len(embedding)} dimensional embedding")

# Test 2: Database operations
if not self.vec_client:
    print("\n❌ Skipping database tests - client not available")
    return

print(f"\n2. Creating table...")
if not self.create_table():
    print("❌ Table creation failed")
    return

print(f"\n3. Testing vector insertion...")
success = self.insert_vector(
    text=test_text,
    metadata={"source": "test", "category": "ml"},
    doc_id="test_doc_1"
)

if not success:
    print("❌ Vector insertion failed")
    return

print(f"\n4. Testing similarity search...")
results = self.search_similar("data processing", limit=3)

if results:
    print("✅ All tests passed!")
    for i, result in enumerate(results[:2]):
        print(f"    Match {i+1}: {result}")
else:
    print("⚠️ No search results found")

```

```
def main():
```

```

"""Main function to test the VectorStore"""
print("🚀 Starting VectorStore setup and testing...")

try:
    # Check environment variables
    if not os.getenv("GOOGLE_API_KEY"):
        print("❌ GOOGLE_API_KEY environment variable not set")
        print("Set it with: export GOOGLE_API_KEY='your_api_key'")
        return

    # Initialize and test VectorStore
    vector_store = VectorStore()

    if vector_store.vec_client:
        vector_store.test_functionality()
    else:
        print("❌ VectorStore initialization failed")

except Exception as e:
    print(f"❌ Error in main: {e}")

if __name__ == "__main__":
    main()

```

Step 6: Usage Instructions

1. **Update database credentials** in the code or set environment variable:

bash

```
export DATABASE_URL="postgresql://vector_user:your_secure_password@localhost:5432/vector_db"
```



2. **Set your Google API key:**

bash

```
export GOOGLE_API_KEY="your_actual_gemini_api_key"
```

3. **Run the script:**

bash

```
python vector_store_fixed.py
```

Common Issues and Solutions

Issue 1: "relation does not exist"

Solution: Run the table creation step in the code

Issue 2: "permission denied"

Solution: Check user privileges in PostgreSQL

Issue 3: "connection refused"

Solution: Ensure PostgreSQL is running:

```
bash
sudo systemctl status postgresql
sudo systemctl start postgresql
```

Issue 4: Wrong database URL format

Solution: Use exact format:

```
postgresql://username:password@hostname:port/database_name
```

This complete setup should resolve all the issues you're experiencing!